

CS525: Advanced Database Organization

Notes 7: Recovery and Concurrency Control Part I: Failure and Recovery

Gerald Balekaki

Department of Computer Science

Illinois Institute of Technology

gbalekaki@iit.edu

July 31st 2023

Slides: adapted from courses taught by [Shun Cheung, Emory University](#), [Andy Pavlo, Carnegie Mellon University](#), [Jennifer Welch, Texas A&M](#) & Introduction to Database Systems by ITL Education Solutions Limited

Concurrency and Recovery

- DBMS should enable reestablish correctness of data in the presence of failures
 - System should restore a correct state after failure (recovery)
- DBMS should enable multiple clients to access the database concurrently
 - This can lead to problems with correctness of data because of interleaving of operations from different clients
 - System should ensure correctness (concurrency control)

Crash Recovery

- Recovery algorithms are techniques to ensure database consistency, transaction atomicity, and durability despite failures.
- Every recovery algorithm has two parts:
 - Actions during normal transaction processing to ensure that the DBMS can recover from a failure.
 - Actions after a failure to recover the database to a state that ensures atomicity, consistency, and durability

Crash Recovery

- DBMS is divided into different components based on the underlying storage device.
- We must also classify the different types of failures that the DBMS needs to handle.

Failure Classification

- **Type #1: Transaction Failures**
- **Type #2: System Failures**
- **Type #3: Storage Media Failures**

- **Logical Errors**

- A transaction cannot complete due to some internal error condition (e.g., integrity constraint violation).

- **Internal State Errors**

- The DBMS must terminate an active transaction due to an error condition (e.g., deadlock)

System Failures

- **Software Failure**

- There is a problem with the DBMS implementation (e.g., uncaught divide-by-zero exception) and the system has to halt

- **Hardware Failure**

- The computer hosting the DBMS crashes (e.g., power plug gets pulled).
- **Fail-stop Assumption**
 - We assume that non-volatile storage contents are not corrupted by system crash.

- **Non-Repairable Hardware Failure**

- A head crash or similar disk failure destroys all or parts of nonvolatile storage.
 - Destruction is assumed to be detectable.
- No DBMS can recover from this. Database must be restored from archived version

Failure Modes: System failure example

- Example: Transaction: transfer \$100 from account A to account B

```
1 READ A
2 A.balance = A.balance - 100
3 WRITE A
4 **system fails here**
5 READ B
6 B.balance = B.balance + 100
7 WRITE B
```

- Then A would lose his \$100 !!!
- This problem is solved by logging
- Transaction needs to be executed correctly

Modeling consistency of a database

- **Database element:** the unit of data accessed by the database system
 - Abstraction that will come in handy when talking about concurrency control and recovery
- **Database:** a collection of database elements
- **Note:**
 - Different DBMS uses different notion for database element
 - Possible units:
 - A relation
 - A disk block
 - A tuple in a relation

Modeling consistency of a database

- Database state: the collection of values of all database elements in the database
- Database state can be changed by changing one or more of the database elements in the database
- A database state can be
 - Consistent: satisfy all constraints of the database schema and implicit constraints
 - Inconsistent


Modeling consistency of a database

- **Transaction:** a sequence of changes to one or more database elements
- **Example:** Transaction: transfer \$100 from account A to account B

```
1 READ A
2 A.balance = A.balance - 100
3 WRITE A
4 READ B
5 B.balance = B.balance + 100
6 WRITE B
```

Modeling consistency of a database

- A more precise definition of transaction:
- **Transaction:** a sequence of changes to one or more database elements
- When all changes in a transaction are made to the database state:
 - The resulting database state is a **consistent state** (if the initial state is consistent)




Causes of inconsistent database states

- There are 2 causes of inconsistent database states
 1. System failure
 2. Concurrent execution

How a system failure can result in an inconsistent DB state

- Consider the following transaction transfer \$100 from A to B

```
1 READ A
2 A.balance = A.balance - 100
3 WRITE A
4 READ B
5 B.balance = B.balance + 100
6 WRITE B
```




- There are 2 possible consistent states
 - Consistent state 1: A B
 - Consistent state 2: A-100 B+100

How a system failure can result in an inconsistent DB state

- Consider the database state that result from the following system failure

```
1 READ A
2 A.balance = A.balance - 100
3 WRITE A
4 **system fails here**
5 READ B
6 B.balance = B.balance + 100
7 WRITE B
```




- The resulting database state is: Database state = A-100 B
- Not one of the 2 possible consistent states


How concurrent execution can cause inconsistent states

- Consider the following 2 transactions

T1= Deposit \$1 to A




T2= Deposit \$2 to A




- The possible consistent database states for executing T1 and T2 are:

- Case 1: T1 before T2



- Case 2: T2 before T1



How concurrent execution can cause inconsistent states

- Consider the following concurrent execution of T1 and T2
- T1= Deposit \$1 to A, T2= Deposit \$2 to A

T1	T2
READ A	
	READ A
$A = A + 1$	$A = A + 2$
WRITE A	
	WRITE A

initially: $A = 10$


($A = 11$)

($A = 12$)

Writes 11 to A

Writes 12 to A

- Final database state: $A = 12$ ($= A + 2$)



Correctness “theory” of database transactions

- Assuming that the database is in a **consistent state**
- Then, a transaction will transform the database into a (another) consistent state if:
 - There are no system failures
 - There are no other transactions executing in the database system

Transactions

- **Transaction**: the (smallest) unit of execution of database operations (updates)
- Unit = whole thing, indivisible
- A **transaction** is:
 - executed completely or
 - nothing from the **transaction** is executed

Transactions in SQL

- A new txn starts with the **BEGIN** command.
- The txn stops with either **COMMIT** or **ABORT**
 - If commit, the DBMS either saves all the txn's changes or aborts it.
 - If abort, all changes are undone so that it's like as if the txn never executed at all.
- Abort can be either self-inflicted or caused by the DBMS.

Notation for a transaction

```
begin transaction
....
.... operations performed by the transaction
.... e.g.: read, compute, write
....
commit    // success
```

• or

```
begin transaction
....
.... operations performed by the transaction
.... e.g.: read, compute, write
....
abort     // failure
```

Notation for a transaction: Result

- All operations between

```
begin transaction
....
....  ALL operations executed
....
commit
```

will be executed

- None of the operations between

```
begin transaction
....
....  NO operations executed
....
abort
```

has been executed

Correctness Criteria: ACID

- A **transaction** (must) have the following properties
 - Atomicity
 - Either all operations of the transaction are properly reflected in the database or none are (i.e., either all or no commands of transaction are executed)
 - Consistency
 - Execution of a transaction in isolation preserves the consistency of the database.
 - Isolation
 - If two **transactions** are executing concurrently, each **transaction** will see the database as if the **transaction** was executing **sequentially** (in isolation) (i.e., transactions are running isolated from each other)
 - Durability
 - After a **transaction** completes successfully, the changes it has made to the database persist (permanent), even if there are system failures (i.e., modifications of transactions are never lost)